

**MIS 382N: Business Data Science | Fall 2019**  
**Homework One**

**Rutwick Bhawsar | Shivang Arya**

**Libraries Imported**

```
import numpy as np
import matplotlib.pyplot as plt
import math
import pandas as pd
from sklearn.decomposition import PCA
```

1. Create 1000 samples from a Gaussian distribution with mean -10 and standard deviation 5. Create another 1000 samples from another independent Gaussian with mean 10 and standard deviation 5.
  - a. Take the sum of 2 these Gaussians by adding the two sets of 1000 points, point by point, and plot the histogram of the resulting 1000 points. What do you observe?
  - b. Estimate the mean and the variance of the sum.

**Answer**

Creating a NumPy array of 1000 samples with mean -10 and std. deviation 5

```
mu = -10
sigma = 5
arr1 = np.random.normal(mu, sigma, 1000)
arr2 = np.random.normal(-mu, sigma, 1000)
```

Adding the two NumPy arrays and getting the mean and standard deviation for the purpose of plotting


```
arr3 = np.add(arr1, arr2)
mu = arr3.mean()
sigma = arr3.std()
```

**Plotting Histogram**

```
count, bins, ignored = plt.hist(arr3, 40, normed=True)
plt.plot(bins, 1/(sigma * np.sqrt(2 * np.pi)) * np.exp(-(bins - mu)**2 / (2 * sigma**2)), linewidth=2, color='r')
plt.show()
```

Calculating and showing mean and variance

```
print("Mean : ", arr3.mean())
print("Variance : ", arr3.var())
```

**Output**

Mean : -0.2953283731254982  
 Variance : 44.44237249218313

2. Central Limit Theorem. Let  $X_i$  be an iid Bernoulli random variable with value  $\{-1,1\}$ . Look at the random variable.  $Z_n = \frac{1}{\sqrt{n}} \sum X_i$ .  
 By taking 1000 draws from  $Z_n$ , plot its histogram. Check that for small  $n$  (say, 5-10)  $Z_n$  does not look that much like a Gaussian, but when  $n$  is bigger (already by the time  $n = 30$  or  $50$ ) it looks much more like a Gaussian. Check also for much bigger  $n$ :  $n = 250$ , to see that at this point, one can really see the bell curve.

**Answer**

Defining a function which returns the value of  $Z_n$ .

```
def iid(n):
    temparr = [-1,1]
    arr1=np.random.choice(temparr,size = n)
    summ = np.sum(arr1)
    rootn = math.sqrt(n)
    z = summ/rootn
    return z
arr = np.array([])
```

A for loop to take 1000 draws and 1000 values of  $Z_n$ . Then appending the value to an array.


```
for i in range(1000):  
    zt=iid(5)  
    arr = np.append(arr,zt)
```

Plotting the array.


```
plt.hist(arr,density="true",bins=100)  
plt.show()
```



### Output

1. Output when  $n = 5$



2. Output when  $n = 45$



3. Output when  $n = 250$ 4. Output when  $n = 5000$ 

3. Estimate the mean and standard deviation from 1 dimensional data: generate 25,000 samples from a Gaussian distribution with mean 0 and standard deviation 5. Then estimate the mean and standard deviation of this Gaussian using elementary numpy commands, i.e., addition, multiplication, division (do not use a command that takes data and returns the mean or standard deviation).

**Answer**

Creating a NumPy array of 25000 samples with mean 0 and std. deviation 5.

```
mu = 0
sigma = 5
n = 25000
array = np.random.normal(mu, sigma, n)
```

Calculating the mean by adding all the elements in the array and dividing by total number of elements.

```
mean = np.sum(array)/n
print("Mean : ",mean)
```

Printing the mean using NumPy functions for verification

```
print("Mean with function for verification : ", np.mean(array))
```

Calculating the std. deviation by using the formula

$$s = \sqrt{\frac{\sum_{i=1}^n (x_i - \bar{x})^2}{n-1}}$$

```
stddev = np.sqrt(np.sum(np.subtract(array,mean)**2)/n)
print("\n\nStd. Deviation : ",stddev)
```

Printing the std. deviation using NumPy functions for verification

```
print("Std. Deviation with function for verification : ",np.std(array))
```


Plotting the histogram

```
count, bins, ignored = plt.hist(array, 30, normed=True)
plt.plot(bins, 1/(sigma * np.sqrt(2 * np.pi)) * np.exp( - (bins - mu)**2 / (2 * sigma**2) ), linewidth=2,
color='black')
plt.show()
```

## Output

```
Mean : 0.08084793674503259
Mean with function for verification : 0.08084793674503259
```

```
Std. Deviation : 4.9977962497816675
Std. Deviation with function for verification : 4.9977962497816675
```



4. Estimate the mean and covariance matrix for multi-dimensional data: generate 10,000 samples of 2 dimensional data from the Gaussian distribution

$$\begin{pmatrix} X_i \\ Y_i \end{pmatrix} \sim \mathcal{N} \left( \begin{pmatrix} -5 \\ 5 \end{pmatrix}, \begin{pmatrix} 20 & .8 \\ .8 & 30 \end{pmatrix} \right)$$

Then, estimate the mean and covariance matrix for this multi-dimensional data using elementary numpy commands, i.e., addition, multiplication, division (do not use a command that takes data and returns the mean or standard deviation). (Recall: this notation is a compact way of specifying the distribution of two Gaussian random variables that are correlated. Specially, the notation indicates that X has mean -5 and 1 variance 20, Y has mean 5 and variance 30, and the covariance of X and Y is 0.8, i.e.,  $E[(X + 5)(Y - 5)] = 0.8$ .)

### **Answer**

Creating two NumPy array of 10000 samples. One with mean -5 and std. deviation = square root of  $\sqrt{20}$ . Another with mean 5 and std. deviation is  $\sqrt{30}$  (*Std. Dev =  $\sqrt{\text{Variance}}$* )

```
mu1 = -5
sigma1 = math.sqrt(20)
mu2 = 5
sigma2 = math.sqrt(30)
n = 10000
ar1 = np.random.normal(mu1,sigma1,n)
ar2 = np.random.normal(mu2,sigma2,n)
```

Taking the mean of both the arrays. Used for covariance matrix.

```
mean1 = np.sum(ar1)/n
mean2 = np.sum(ar2)/n
```

Creating a 2D array with both the previous arrays.

```
arr = np.array([ar1,ar2])
```

Taking the mean and verifying it

```
mean = (np.sum(arr)/(n*2))
print("Mean : ",mean)
print("Mean with function for verification : ",np.mean(arr))
```

**Calculating  $Cov_{xy}$ ,  $Cov_{xx}$ ,  $Cov_{yy}$** 

```
covariancexy = (np.sum(np.subtract(ar1,mean1)*np.subtract(ar2,mean2))/n)
covariancexx = np.sum(np.subtract(ar1,mean1)**2)/n
covarianceyy = np.sum(np.subtract(ar2,mean2)**2)/n
```

**Putting the values in the Covariance Matrix**

```
covarr = np.array([[covariancexx,covariancexy],[covariancexy,covarianceyy]])
```

**Printing and Verifying the matrix**

```
print("\nCovariance Array: \n", covarr)
print("\nCovariance with function for verification : \n",np.cov(arr))
```

**Output**

```
Mean : 0.010484252462945733
Mean with function for verification : 0.010484252462945733

Covariance Array:
[[20.20768136 -0.30684038]
 [-0.30684038 30.06494902]]

Covariance with function for verification :
[[20.20970233 -0.30687106]
 [-0.30687106 30.06795581]]
```

**5. Download from Canvas/Files the dataset PatientData.csv.**

Each row is a patient and the last column is the condition that the patient has. Do data exploration using Pandas and other visualization tools to understand what you can about the dataset. For example:

- How many patients and how many features are there?
- What is the meaning of the first 4 features? See if you can understand what they mean.
- Are there missing values? Replace them with the average of the corresponding feature column
- How could you test which features strongly influence the patient condition and which do not? List what you think are the three most important features.

**Answer**

Read the .csv file in Pandas dataframe df and store a backup of in in df2.

```
df = pd.read_csv('PatientData.csv',header=None)
df2 = df
```

a. Print the shape of the dataframe to identify the number of rows and column

```
print(df.shape)
```

**Output**

```
(452, 280)
```

The total Rows = 452

The total Columns = 280

b. We followed the following steps to determine the meaning of the first 4 columns.

i. First we extracted the first 4 column from the entire dataframe. Then described it.

```
df = df[[0,1,2,3]].describe()
```

	0	1	2	3
count	452.000000	452.000000	452.000000	452.000000
mean	46.471239	0.550885	166.188053	68.170354
std	16.466631	0.497955	37.170340	16.590803
min	0.000000	0.000000	105.000000	6.000000
25%	36.000000	0.000000	160.000000	59.000000
50%	47.000000	1.000000	164.000000	68.000000
75%	58.000000	1.000000	170.000000	79.000000
max	83.000000	1.000000	780.000000	176.000000

ii. We sorted the data to find the outliers.

```
df = df.sort_values(2,ascending = False)
```

iii. We found that the top 4 rows in the 3<sup>rd</sup> column were noise and so we eliminated them. Then we described the table again.

```
df = df.iloc[3:]
df.describe()
```



	0	1	2	3
count	449.000000	449.000000	449.000000	449.000000
mean	46.697105	0.552339	163.783964	68.400891
std	16.225862	0.497808	10.350126	16.131701
min	1.000000	0.000000	105.000000	10.000000
25%	36.000000	0.000000	160.000000	59.000000
50%	47.000000	1.000000	164.000000	68.000000
75%	58.000000	1.000000	170.000000	79.000000
max	83.000000	1.000000	190.000000	176.000000

- iv. Based on the raw data, mean, std. deviation and minimum and maximum values we deduced the following.

**First Column: Age.** Because the max is 83 and min is 1 and the average is 46.6.

**Second Column: Gender.** Because the columns had only 0s and 1s so we realized that it was a categorical variable. Since the mean was close to 0.5 we assumed that there were approximately equal male and females.

**Third Column: Height in cms.** Because avg. was 163, min was 105 and max was 190 which is what we figured to be the average heights.

**Fourth Column: Weight in kgs.** Because firstly the units are in S.I. Secondly looking at min is 10kgs, maximum is 176kgs which is slightly high but it's believable.

- v. To confirm our deduction found correlation amongst these column since in reality age, height and weight are related.

Finding the correlation between what we think is age and height in cms

```
print("\n",df[0].corr(df[2]))
```

Finding the correlation between what we think is age and weight in kgs

```
print(df[0].corr(df[3]))
```

Finding the correlation between what we think is height and weight

```
print(df[2].corr(df[3]))
```

Our assumptions were confirmed since we found a positive correlation amongst all of them.

## Output

The correlation between what we think is age and height 0.24850302559483486

The correlation between what we think is age and weight 0.35490347920428017

The correlation between what we think is height and weight 0.5730169619733592

vi. We first grouped the data on the basis on gender and got the means.

We then inferred that 0 is Male and 1 is Female. Since Avg. Height and Avg. Weight of 0s was higher than those of 1.

```
df_gender=df[[0,1,2,3]].groupby(1).mean()
df_gender.head()
```

### Output

	0	2	3
1	Age	Height	Weight
0	47.646766	169.049751	73.000000
1	45.774194	159.516129	64.693548

c. We observed the dataset and observed that the missing values were denoted by '?' instead of NaN. First we replaced the '?' with NaN

```
df = df[:].replace('?', np.nan)
```

Then we figured out which columns had missing values

```
print(df.columns[df.isna().any()].tolist())
```

```
[10, 11, 12, 13, 14]
```

After describing these columns we realized that the mean wasn't getting calculated. This was because the values in the columns were stored as strings. So we explicitly converted them to float and then described them.

```
df.iloc[:,10:15]=df.iloc[:,10:15].astype(float)
print(df.iloc[:,10:15].describe())
```

	10	11	12	13	14
count	444.000000	430.000000	451.000000	76.000000	451.000000
mean	36.150901	48.913953	36.716186	-13.592105	74.463415
std	57.858255	29.346409	36.020725	127.220248	13.870684
min	-177.000000	-170.000000	-135.000000	-179.000000	44.000000
25%	14.000000	41.000000	12.000000	-124.500000	65.000000
50%	41.000000	56.000000	40.000000	-50.500000	72.000000
75%	63.250000	65.000000	62.000000	117.250000	81.000000
max	179.000000	176.000000	166.000000	178.000000	163.000000


We replaced NaN with the mean of the corresponding column. To confirm the replacement, we described it again, saw that the mean was unchanged and checked if there are any columns with missing values and the answer was no.

```
df = df.fillna(df.mean(axis = 0, skipna = True))
print("\n",df.iloc[:,10:15].describe())
print("\n",df.columns[df.isna().any()].tolist())
```

	10	11	12	13	14
count	452.000000	452.000000	452.000000	452.000000	452.000000
mean	36.150901	48.913953	36.716186	-13.592105	74.463415
std	57.342803	28.621694	35.980768	51.879836	13.855298
min	-177.000000	-170.000000	-135.000000	-179.000000	44.000000
25%	14.000000	41.000000	12.000000	-13.592105	65.000000
50%	41.000000	54.500000	40.000000	-13.592105	72.000000
75%	63.000000	64.000000	62.000000	-13.592105	81.000000
max	179.000000	176.000000	166.000000	178.000000	163.000000

[ ]

d. For finding the most important features in our dataset we used Principal Component Analysis. We extracted 7 components based on the Proportion of Variance explained graph.



```
pca = PCA(n_components = 7)
pca
pca.fit(df)
```

Proportion of Variance explained by each principal component.

```
print('Explained variation per principal component: {}'.format(pca.explained_variance_ratio_))
```

```
Explained variation per principal component: [0.15476467 0.10033137 0.0932275  0.08503253 0.05277753 0.04974104
0.0462441 ]
```

We found out the features which have the maximum influence on our 7 principal components.

Features	PC1	Features	PC2	Features	PC7
9	0.444096	248	0.412921	7	0.351471
12	0.304351	247	0.359577	56	0.271770
258	0.230049	238	0.280944	5	0.245042
13	0.219639	258	0.275484	4	0.227163
248	0.174147	9	0.233061	6	0.218120
40	0.172991	237	0.231513	77	0.201978
187	0.170932	257	0.227793	29	0.192762
268	0.168229	12	0.169756	92	0.190674
257	0.167631	87	0.168455	17	0.187015
64	0.165991	99	0.160757	149	0.177752

The last column in the dataset is the dependent variable. This is stored in another dataset.

```
dfultimate=df.iloc[:, -1]
```

We fit all the rows in our dataset to the 7 principle components.

```
print(fitted)
principal_new = pd.DataFrame(data = fitted, columns = ['principal component 1', 'principal component 2',
'principal component 3', 'principal component 4', 'principal component 5', 'principal component 6', 'principal
component 7'])
principal_new
```

	principal component 1	principal component 2	principal component 3	principal component 4	principal component 5	principal component 6	principal component 7
0	32.712062	-35.821108	45.838801	15.276055	61.508009	-6.511624	-32.633378
1	-9.811479	-9.619186	24.045219	25.727460	-7.549748	16.378937	-41.057002
2	-72.619499	23.043276	-98.006948	-72.093470	45.683185	17.372803	128.273010
3	10.230763	-48.173107	29.142867	50.251214	72.729612	-0.653305	19.941414
4	47.995382	-28.980950	29.100665	-12.794413	41.062773	-8.905969	-21.810897
5	-104.013522	169.411578	-94.149530	31.509676	-3.947101	13.798762	48.412344
6	-27.127933	49.560860	-27.407973	-9.834545	-44.549169	4.054703	-14.595423
7	-15.972302	48.149912	23.891734	-49.190753	-122.948828	47.449451	49.975615
8	-23.713266	-7.602419	-73.910951	-1.341418	-50.433647	86.614372	-12.394695
9	-75.624734	31.027837	-35.332997	-34.635360	-16.708726	4.698701	48.091979

We are finding the regression equation for all 7 of our principle components with respect to the dependent variable i.e. Last Column.

```
from sklearn import linear_model
regr = linear_model.LinearRegression()
regr.fit(principal_new, dfultimate)
predicted = regr.predict(principal_new)
mse = np.mean((predicted-dfultimate)**2)
print (regr.intercept_, regr.coef_,mse)
print(regr.score(principal_new, dfultimate))
```

```
3.880530973451327 [ 0.00131832  0.00610902 -0.00407609 -0.00123462 -0.00458422 -0.00024039
 0.03308703] 16.844482251106612
0.1308107001739479
```

What we understand here is that the 7th Principal Component has the most effect on the independent variable by a large extent (coefficient of regression is more than 10 times higher than the other principal components).

Hence, we choose the three features that affect Principal Component 7 the most:

1. Feature 7 in Python (8th column in CSV)
2. Feature 56 in Python (57th column in CSV)
3. Feature 5 in Python (6th column in CSV)

## Part II – Written Question on next page.

P.T.O



	$X=0$	$X=1$	Total
$Y=0$	$\frac{1}{4}$	$\frac{1}{4}$	$\frac{1}{2}$
$Y=1$	$\frac{1}{6}$	$\frac{1}{3}$	$\frac{1}{2}$
Total	$\frac{5}{12}$	$\frac{7}{12}$	

(a)  $\rightarrow P(X=1) = ?$

From above table, it is clear that

$$\underline{\underline{P(X=1) = \frac{7}{12}}}$$

(b)  $\rightarrow P(X=1 | Y=1) = ?$

From above table,

$$\underline{\underline{P(X=1 | Y=1) = \frac{1}{3} \times \frac{1}{(\frac{1}{3} + \frac{1}{6})} = \frac{2}{3}}}$$

(c)  $\rightarrow$  Variance of random variable  $X$ ?

According to  $X$ 's probability distribution table, in 12 chances there is a high probability of  $X$  being 1 7 times and 0 for the remaining 5.

Therefore,

$$\text{mean}(x) = 0\left(\frac{5}{12}\right) + 1\left(\frac{7}{12}\right)$$

$$\Rightarrow \mu_x = \frac{7}{12}$$

$x$	$(x - \mu)^2$
1	$\frac{25}{144}$
1	$\frac{25}{144}$
1	$\frac{25}{144}$
1	$\frac{25}{144}$
1	$\frac{25}{144}$
1	$\frac{25}{144}$
1	$\frac{25}{144}$
0	$\frac{49}{144}$
0	$\frac{49}{144}$
0	$\frac{49}{144}$
0	$\frac{49}{144}$
0	$\frac{49}{144}$

$$E(x - \mu)^2 = \frac{(25 \times 7)}{144} + \frac{(49 \times 5)}{144}$$

$$= \frac{175}{144} + \frac{245}{144} = \frac{420}{144}$$

$$\text{Variance } (\sigma_x^2) = \frac{\sum (x - \mu)^2}{n}$$

$$= \frac{420}{144} \times \frac{1}{12}$$

$$= \frac{35}{144} \text{ or } 0.2430$$

(d) → variance of  $x$  given that  $y=1$  ?  
 $(\sigma_{x|y=1}^2)$

Conditional distribution of  $x$  given that  $y=1$  ~~(scribbled out)~~

	$x=0$	$x=1$	Total
$y=1$	$\frac{1/6}{1/6+1/3} = \frac{1}{3}$	$\frac{1/3}{1/6+1/3} = \frac{2}{3}$	1

mean of  $x$  given  $y=1$

$$(\mu_{x|y=1}) = \frac{1}{3}(0) + \frac{2}{3}(1)$$

$$= \frac{2}{3}$$

$x$	$(x - \mu)^2$	$\sigma_{x y=1}^2 = \frac{\sum (x - \mu)^2}{n}$
0	4/9	
1	1/9	$\Rightarrow \sigma_{x y=1}^2 = \frac{2}{3} \times \frac{1}{3} = \frac{2}{9}$
1	1/9	

$$\Rightarrow \sigma_{x|y=1}^2 = \frac{2}{3} \times \frac{1}{3} = \frac{2}{9}$$



$$(e) \rightarrow E [x^3 + x^2 + 3y^7 \mid y=1] = ?$$

Given that  $y=1$ , we can use the same conditional distribution as the previous part of the question.

$x$	$x^2$	$x^3$	$3y^7$	$x^3 + x^2 + 3y^7$
0	0	0	3	3
1	1	1	3	5
1	1	1	3	5

$$E [x^3 + x^2 + 3y^7 \mid y=1] = \frac{\sum (x^3 + x^2 + 3y^7)}{n}$$

$$= \frac{13}{3}$$

$$\underline{\underline{= \frac{13}{3}}}$$

2.

$$v_1 = [1, 1, 1]$$

$$v_2 = [1, 0, 0]$$

$v_1 \times v_2$  will give a vector perpendicular to both of them.

Therefore we can project points on the subspace covered by  $v_1$  &  $v_2$  by subtracting the perpendicular vector from the points.

$$v_1 \times v_2 = \begin{vmatrix} \hat{i} & \hat{j} & \hat{k} \\ 1 & 1 & 1 \\ 1 & 0 & 0 \end{vmatrix}$$

$$= \hat{i}(0) - \hat{j}(-1) + \hat{k}(-1)$$

$$\vec{v}_3 = 0\hat{i} + 1\hat{j} - 1\hat{k}$$

Projection of  $P_1 [3, 3, 3]$  on  $\mathbb{R}^3$

$$= [3, 3, 3] - \frac{P_1 \cdot [0, 1, -1]}{(\sqrt{2})^2} [0, 1, -1]$$

$$= [3, 3, 3] - [0, 0, 0]$$

$$= [3, 3, 3]$$

Projection of  $P_2$  on  $R^3$   
 $[1, 2, 3]$

$$= [1, 2, 3] - \frac{[1, 2, 3] \cdot [0, 1, -1]}{(\sqrt{2})^2} [0, 1, -1]$$

$$= [1, 2, 3] - \frac{0 + 2 - 3}{2} [0, 1, -1]$$

$$= [1, 2, 3] - [0, \frac{-1}{2}, \frac{1}{2}]$$

$$= [1, 2.5, 2.5]$$

Projection of  $P_3$  on  $R^3$   
 $[0, 0, 1]$

$$= [0, 0, 1] - \frac{[0, 0, 1] \cdot [0, 1, -1]}{(\sqrt{2})^2} [0, 1, -1]$$

$$= [0, 0, 1] - \frac{0 + 0 - 1}{2} [0, 1, -1]$$

$$= [0, 0, 1] - [0, \frac{-1}{2}, \frac{1}{2}]$$

$$= [0, \frac{1}{2}, \frac{1}{2}]$$

3 → Given  $P(X) = \frac{2}{3}$ ,  $n = 100$

$X \rightarrow$  event of getting a Heads

$n \rightarrow$  number of coin tosses

$x \rightarrow$  outcome we want

To calculate:  $P(X \leq x)$   
where  $x = 50$

Using the Binomial Distribution,


$$P(X \leq 50) = {}^{100}C_0 \left(\frac{2}{3}\right)^0 \left(\frac{1}{3}\right)^{100} + {}^{100}C_1 \left(\frac{2}{3}\right)^1 \left(\frac{1}{3}\right)^{99} + \dots + {}^{100}C_{50} \left(\frac{2}{3}\right)^{50} \left(\frac{1}{3}\right)^{50}$$

$$= 0.0004182$$

$$\approx 0.00042$$

---





Confirmed Question 3 using Binomial Distribution Calculator online.